

# Logistic Regression -1

```
In [1]: import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv("1_ionosphere.csv")  
df
```

```
Out[3]:
```

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...	-0.51171	C
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.26569	-1
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.40220	1
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.90695	1
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.65158	1
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.01535	-1
...	...	...	...	...	...	...	...	...	...	...	...	...	...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-0.04202	1
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	0.01361	1
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	0.03193	1
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...	-0.02099	1
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...	-0.15114	1

350 rows × 35 columns

```
In [4]: feature_matrix=df.iloc[:,0:34]  
target_vector=df.iloc[:, -1]
```

```
In [5]: feature_matrix.shape
```

```
Out[5]: (350, 34)
```

```
In [6]: target_vector.shape
```

```
Out[6]: (350,)
```

```
In [7]: from sklearn.preprocessing import StandardScaler
```

```
In [10]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [11]: logr=LogisticRegression()
logr.fit(fs,target_vector)
```

```
Out[11]: LogisticRegression()
```

```
In [12]: observation=[[1.4,2.3,-5.9,11,12,13,14,15,16,17,1,2,3,4,5,6,7,8,9,10,21,22,23,24,25,26,
```

```
In [15]: prediction=logr.predict(observation)
print(prediction)
```

```
['g']
```

```
In [16]: logr.classes_
```

```
Out[16]: array(['b', 'g'], dtype=object)
```

```
In [18]: logr.predict_proba(observation)[0][0]
```

```
Out[18]: 0.0
```

```
In [19]: logr.predict_proba(observation)[0][1]
```

```
Out[19]: 1.0
```

## Logistic Regression - 2

```
In [20]: import re
from sklearn.datasets import load_digits
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

```
In [21]: digits=load_digits()
digits
```

```
Out[21]: {'data': array([[ 0.,  0.,  5., ...,  0.,  0.,  0.],
 [ 0.,  0.,  0., ..., 10.,  0.,  0.],
 [ 0.,  0.,  0., ..., 16.,  9.,  0.],
```

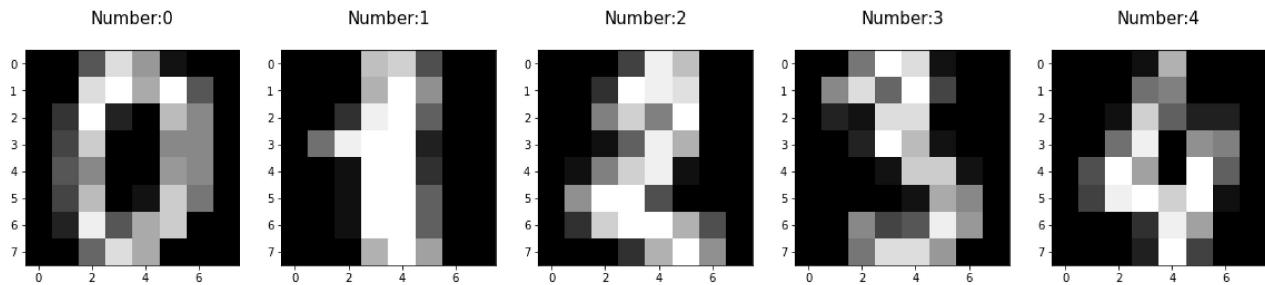
```
...,
[ 0.,  0.,  1., ...,  6.,  0.,  0.],
[ 0.,  0.,  2., ..., 12.,  0.,  0.],
[ 0.,  0., 10., ..., 12.,  1.,  0.]]),
'target': array([0, 1, 2, ..., 8, 9, 8]),
'frame': None,
'feature_names': ['pixel_0_0',
'pixel_0_1',
'pixel_0_2',
'pixel_0_3',
'pixel_0_4',
'pixel_0_5',
'pixel_0_6',
'pixel_0_7',
'pixel_1_0',
'pixel_1_1',
'pixel_1_2',
'pixel_1_3',
'pixel_1_4',
'pixel_1_5',
'pixel_1_6',
'pixel_1_7',
'pixel_2_0',
'pixel_2_1',
'pixel_2_2',
'pixel_2_3',
'pixel_2_4',
'pixel_2_5',
'pixel_2_6',
'pixel_2_7',
'pixel_3_0',
'pixel_3_1',
'pixel_3_2',
'pixel_3_3',
'pixel_3_4',
'pixel_3_5',
'pixel_3_6',
'pixel_3_7',
'pixel_4_0',
'pixel_4_1',
'pixel_4_2',
'pixel_4_3',
'pixel_4_4',
'pixel_4_5',
'pixel_4_6',
'pixel_4_7',
'pixel_5_0',
'pixel_5_1',
'pixel_5_2',
'pixel_5_3',
'pixel_5_4',
'pixel_5_5',
'pixel_5_6',
'pixel_5_7',
'pixel_6_0',
'pixel_6_1',
'pixel_6_2',
'pixel_6_3',
'pixel_6_4',
'pixel_6_5',
'pixel_6_6',
'pixel_6_7',
'pixel_7_0',
'pixel_7_1',
'pixel_7_2',
```



g set and different 13\nto the test set. 32x32 bitmaps are divided into nonoverlapping blocks of\nto 4x4 and the number of on pixels are counted in each block. This generates\nan input matrix of 8x8 where each element is an integer in the range\nto 16. This reduces d imensionality and gives invariance to small\ndistortions.\n\nFor info on NIST preprocess ing routines, see M. D. Garris, J. L. Blue, G.\nT. Candela, D. L. Dimmick, J. Geist, P. J. Grother, S. A. Janet, and C.\nL. Wilson, NIST Form-Based Handprint Recognition Syste m, NISTIR 5469,\n1994.\n.. topic:: References\n - C. Kaynak (1995) Methods of Combi ning Multiple Classifiers and Their\n Applications to Handwritten Digit Recognition, MSc Thesis, Institute of\n Graduate Studies in Science and Engineering, Bogazici Univ ersity.\n - E. Alpaydin, C. Kaynak (1998) Cascading Classifiers, Kybernetika.\n - Ken Tang and Ponnuthurai N. Suganthan and Xi Yao and A. Kai Qin.\n Linear dimensionalityr eduction using relevance weighted LDA. School of\n Electrical and Electronic Engineer ing Nanyang Technological University.\n 2005.\n - Claudio Gentile. A New Approximate Maximal Margin Classification\n Algorithm. NIPS. 2000.\n"}

In [26]:

```
plt.figure(figsize=(20,4))
for index,(image,label) in enumerate(zip(digits.data[0:5],digits.target[0:5])):
    plt.subplot(1,5,index+1)
    plt.imshow(np.reshape(image,(8,8)),cmap=plt.cm.gray)
    plt.title('Number:%i\n'%label,fontsize=15)
```



In [27]:

```
x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,test_size=0.30)
```

In [28]:

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(1257, 64)
(540, 64)
(1257,)
(540,)
```

In [30]:

```
logre=LogisticRegression(max_iter=10000)
logre.fit(x_train,y_train)
```

Out[30]: LogisticRegression(max\_iter=10000)

In [31]:

```
print(logre.predict(x_test))
```

```
[4 2 5 5 4 3 6 8 2 6 2 4 7 7 2 4 5 0 1 1 5 9 2 3 3 1 2 0 0 5 6 0 3 4 3 1 7
8 5 5 0 4 6 5 4 5 3 1 2 3 4 3 2 9 7 8 5 5 2 9 7 7 5 4 2 6 1 2 1 8 2 3 6 0
6 3 7 9 9 7 3 8 8 0 8 8 1 9 7 1 8 9 2 8 7 6 5 5 4 6 3 1 5 1 5 6 7 7 7 1 0
0 8 7 4 1 3 1 1 9 1 7 5 3 7 6 5 7 7 3 9 3 7 2 1 1 7 4 3 0 5 8 7 8 1 8 3 2
3 6 6 6 4 4 8 2 0 6 6 3 3 8 9 9 6 1 2 9 6 9 2 7 9 1 0 3 0 7 9 0 2 2 4 9 6
1 0 4 5 1 5 6 0 4 8 3 7 0 4 1 7 2 7 9 4 5 6 5 2 5 4 8 1 2 5 8 1 5 3 6 4 6
6 8 6 3 5 0 4 7 2 0 4 1 8 8 8 2 9 9 1 7 5 9 4 3 3 1 3 0 6 7 5 2 2 6 2 6 8]
```

```

6 1 1 8 9 0 8 2 3 1 7 8 6 2 7 5 2 7 2 9 1 3 5 3 2 9 8 8 6 6 1 5 6 6 9 8 6
7 8 4 6 1 6 3 5 7 6 3 0 2 4 1 8 8 7 1 4 2 7 7 5 0 7 5 4 8 1 3 9 2 4 8 9 1
4 7 3 2 9 0 9 2 4 4 2 5 8 3 6 6 7 8 2 0 8 4 5 7 5 2 7 4 4 7 7 5 9 2 8 6 0
5 7 8 4 9 3 8 3 6 4 2 4 7 4 6 9 2 1 9 9 0 1 6 2 4 0 7 8 4 6 2 9 2 6 0 6 4
9 5 0 3 0 3 7 3 2 4 8 9 2 3 7 4 5 6 3 8 3 1 6 8 1 0 8 6 3 1 9 0 4 7 9 6 3
2 5 4 1 3 8 2 8 5 1 3 4 0 7 9 0 3 4 4 9 4 3 7 5 0 4 9 2 3 9 7 8 1 6 9 1 0
8 4 6 2 9 0 6 1 4 2 2 9 6 5 8 0 0 7 0 1 1 1 0 2 2 4 2 3 4 0 4 6 0 3 9 0 3
3 4 1 5 5 4 2 3 3 7 1 3 6 4 7 6 0 9 2 7 9 0]

```

In [32]: `print(logre.score(x_test,y_test))`

```
0.9592592592592593
```

## Random Forest

In [4]: `import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns`

In [5]: `df=pd.read_csv("1_ionosphere.csv")`

In [8]: `df['g'].value_counts()`

Out[8]: `g 224  
b 126  
Name: g, dtype: int64`

In [9]: `x=df.drop('g',axis=1)  
y=df['g']`

In [10]: `g1={"g":{"g":1,"b":2}}  
df=df.replace(g1)  
df`

Out[10]:

	<b>1</b>	<b>0</b>	<b>0.99539</b>	<b>-0.05889</b>	<b>0.85243</b>	<b>0.02306</b>	<b>0.83398</b>	<b>-0.37708</b>	<b>1.1</b>	<b>0.03760</b>	...	<b>-0.51171</b>	C
<b>0</b>	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.26569	-1
<b>1</b>	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.40220	1
<b>2</b>	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.90695	1
<b>3</b>	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.65158	1
<b>4</b>	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.01535	-1
...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>345</b>	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-0.04202	1
<b>346</b>	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	0.01361	1
<b>347</b>	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	0.03193	1

```
    1   0   0.99539 -0.05889  0.85243  0.02306  0.83398 -0.37708      1.1  0.03760 ... -0.51171 C
348  1   0   0.90608 -0.01657  0.98122 -0.01989  0.95691 -0.03646  0.85746  0.00110 ... -0.02099 I
349  1   0   0.84710  0.13533  0.73638 -0.06151  0.87873  0.08260  0.88928 -0.09139 ... -0.15114 I
```

350 rows × 35 columns

In [11]: `from sklearn.model_selection import train_test_split`

In [12]: `x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)`

In [13]: `from sklearn.ensemble import RandomForestClassifier`

In [14]: `rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)`

Out[14]: `RandomForestClassifier()`

In [24]: `parameters={'max_depth':[1,2,3,4,5],  
 'min_samples_leaf':[5,10,15,20,25],  
 'n_estimators':[10,20,30,40,50]}  
}`

In [25]: `from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)`

Out[25]: `GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param_grid={'max_depth': [1, 2, 3, 4, 5],  
 'min_samples_leaf': [5, 10, 15, 20, 25],  
 'n_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')`

In [26]: `grid_search.best_score_`

Out[26]: `0.9385245901639344`

In [28]: `rfc_best=grid_search.best_estimator_`

In [30]: `from sklearn.tree import plot_tree`

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],fill
```

Out[30]: `[Text(1488.0, 1956.96, '0.85243 <= 0.033\nngini = 0.461\nsamples = 159\nvalue = [88, 156]\n\nclass = No'),`

```

Text(992.0, 1522.080000000002, 'gini = 0.0\nsamples = 33\nvalue = [54, 0]\nclass = Yes'),
Text(1984.0, 1522.080000000002, '-0.51171 <= -0.981\n gini = 0.294\nsamples = 126\nvalue = [34, 156]\nclass = No'),
Text(992.0, 1087.2, '0.59755 <= 0.5\n gini = 0.219\nsamples = 13\nvalue = [14, 2]\nclass = Yes'),
Text(496.0, 652.320000000002, 'gini = 0.48\nsamples = 5\nvalue = [3, 2]\nclass = Yes'),
Text(1488.0, 652.320000000002, 'gini = 0.0\nsamples = 8\nvalue = [11, 0]\nclass = Yes'),
Text(2976.0, 1087.2, '-0.17755 <= -0.749\n gini = 0.203\nsamples = 113\nvalue = [20, 154]\nclass = No'),
Text(2480.0, 652.320000000002, 'gini = 0.278\nsamples = 5\nvalue = [5, 1]\nclass = Yes'),
Text(3472.0, 652.320000000002, '0.59755 <= 0.999\n gini = 0.163\nsamples = 108\nvalue = [15, 153]\nclass = No'),
Text(2976.0, 217.440000000005, 'gini = 0.091\nsamples = 93\nvalue = [7, 139]\nclass = No'),
Text(3968.0, 217.440000000005, 'gini = 0.463\nsamples = 15\nvalue = [8, 14]\nclass = No'))

```

